

Neural MCPF Allocation

Learning the Goal-Allocation Step of Multi-Agent Path Finding

Ofek Yabo | Dayan Badalbaev

Multi-Agent Systems, 2026

Abstract

Multi-agent path finding with multiple goals requires deciding which agent visits which goals before any path can be planned. This allocation sub-problem is a multi-agent traveling salesman problem, and it is the combinatorial part of an otherwise tractable pipeline. The exact RobustMCPF solver handles it with an LKH TSP call, then passes a fixed allocation to conflict-based search (CBS) for collision-free planning. We replace only the allocation step with a learned, size-agnostic transformer. It reads two shortest-path distance matrices, one from agents to goals and one between goals, and outputs a distribution over agents for each goal. Since CBS is shared by both pipelines, any difference in the final plan comes from the allocator alone, and a single set of weights serves any problem size from 2×2 to 50×100 . Across 28 configurations the learned allocator stays within 1–6% of the exact solver’s true execution cost and produces no infeasible allocations. At paper scale it runs $5.9\text{--}9.4\times$ faster while remaining within roughly 20% of optimal. We then ask whether map-specialized allocators beat a single generalist. Five models, one trained on all four benchmark maps and four trained on a single map each, are compared on solution quality and wall-time to determine whether map-aware routing is worthwhile. **[TODO: closing sentence with the headline result, once Exp 16 lands]**

1 Introduction

Fleets of autonomous agents such as warehouse robots, delivery drones and search-and-rescue teams all face the same first decision: given a set of targets and a team of agents, who goes where? The choice matters. A poor allocation sends agents across each other’s paths, inflates total travel, and creates collisions that a planner then has to resolve. It also has to be made before anything else, since path planning only becomes meaningful once each agent knows its goals.

Formally this is multi-agent path finding (MAPF) with the added requirement that goals be assigned. The assignment is not a balanced one-to-one matching. An agent may profitably take zero, one, or several goals, which makes the structure a multi-agent traveling salesman problem (mTSP). The RobustMCPF solver [1] solves it exactly in two stages. An LKH TSP call [3] produces the optimal allocation together with each agent’s visit order, and conflict-based search (CBS) [2] converts that into collision-free paths. The two stages scale very differently. CBS cost depends on how much the agents interfere with each other, whereas the TSP allocation is combinatorial in the number of agents and goals, and it is the stage that degrades first as problems grow.

This work replaces that stage, and only that stage, with a neural network that reads two breadth-first-search distance matrices: D from agents to goals, and G between goals. The downstream ordering and CBS steps are left untouched, so both pipelines end in the same planner and any difference between them can be attributed to the allocator. The trade is exactness for speed. A network trained by imitation will not always reproduce the optimum, and the question this report answers is when accepting that is worthwhile.

Contributions.

- A size-agnostic transformer allocator: one set of weights for any (N, M) , with no positional embeddings and no padding (Section 3).
- Evaluation by true CBS execution cost instead of allocation accuracy, which shows that most disagreements with the solver are cost-equivalent ties (Section 4).

- A design study isolating input representation as the dominant factor, worth more than architecture, capacity and data scale combined (Section 5).
- A characterization of what transfers: map geometry does, problem scale does not (Section 5).
- A scale-limit and specialist-versus-generalist study producing a concrete map-routing recommendation (Section 6).

2 Background and Problem Setting

2.1 Why replace the allocator

The practical case rests on where the cost sits and how often the problem gets solved. In a warehouse or a delivery fleet, allocation is re-solved continuously as orders arrive and agents finish their tasks, and every re-solve pays the combinatorial cost again. An exact allocator is an excellent choice for a single offline instance and a progressively worse one as throughput requirements rise. A learned allocator has the opposite profile. Training is expensive and paid once, while inference is a forward pass whose cost is close to flat in problem size. What is given up is quality, so the report’s job is to measure that loss in the currency that matters, namely the length of the collision-free plan actually executed, and to identify where the exchange pays off.

2.2 Problem formulation

We work in basic MAPF: a 4-connected grid with static obstacles, deterministic movement, no orientation and no action delays. An instance consists of N agents at distinct start cells and M goal cells. Every goal is visited exactly once by exactly one agent, and agents are otherwise unconstrained. The objective is min-sum, minimizing total path length, subject to the paths being collision-free.

That freedom is what makes this an mTSP and not an assignment problem. Figure 1 shows the canonical case of two agents near each other and two distant goals. Splitting the goals costs 14, while giving both to one agent costs 7. The optimum is unbalanced.




Figure 1: Goal allocation is an mTSP, not a balanced assignment. Bundling both goals onto one agent (cost 7) beats splitting them (cost 14). As a result the ground-truth allocation Y has columns that sum to one and rows that are free.

Write $Y \in \{0, 1\}^{N \times M}$ with $Y_{ij} = 1$ when agent i visits goal j . Columns sum to one and rows are unconstrained. The model therefore has to produce a column-wise distribution, giving for each goal a distribution over agents, instead of a doubly-stochastic or permutation-like output. This rules out Sinkhorn-style balanced-assignment layers, which impose a row constraint the problem does not have.

2.3 The exact baseline

RobustMCPF in BasicMAPF mode serves as both ground truth and comparison point. It calls LKH to solve the underlying TSP, which yields k -best allocations along with each agent’s visit order, and then runs CBS to resolve conflicts. We take the cost of the final collision-free plan as

the definitive quality measure. The same solver generates the training labels: every training instance is solved exactly, and its allocation becomes the supervision target.

2.4 Related work

Learning to solve combinatorial problems has a substantial literature. Pointer Networks [9] output permutations of their input, reinforcement-learning variants scaled the idea to routing [10], and attention models became the standard backbone for TSP and VRP [11]; [12] surveys the area. Within MAPF, CBS [2] defines the exact-planning baseline, benchmarks and terminology come from [4, 5], and learned policies such as PRIMAL [8] target path planning, not allocation. Combined target assignment and planning was formalized by [6]. Closest to our final study is algorithm selection for MAPF [7], which asks the same underlying question of whether instance-specific specialization beats one general method. We differ from the routing literature in learning only the allocation and leaving an exact planner downstream, which makes the comparison end-to-end and unambiguous.

3 Method: The Learned Allocator

3.1 Overview

Both pipelines share everything except the allocator (Figure 2). The solver runs LKH followed by CBS. Ours runs a forward pass, then goal-visit ordering, then CBS. Because the planner, the instance and the cost measure are identical, the difference between the two output plans isolates the allocation decision.




Figure 2: Only the allocation box changes. CBS is shared, so any difference in execution cost or wall-time can be attributed to the allocator.

3.2 Input representation

The network never sees the grid. It sees two matrices of BFS shortest-path distances that respect walls: $D \in \mathbb{R}^{N \times M}$ from each agent to each goal, and $G \in \mathbb{R}^{M \times M}$ between goals (Figure 3). Both are normalized to $[0, 1]$ by $(W - 1) + (H - 1)$, the maximum BFS distance on an open $W \times H$ grid.

Using distances instead of coordinates is a deliberate choice. Tour cost is a function of distances, and distances abstract away the map, so the representation transfers across grid sizes and geometries. D tells each agent how far each goal is. G carries the information D cannot express, which is which goals lie near each other and should therefore be bundled onto one agent. Section 5 shows that G is the single most valuable design decision in the project.

3.3 Handling any number of agents and goals

A conventional network fixes its input dimension, which would force one model per (N, M) pair. Ours does not, and the mechanism is worth setting out in full (Figure 4).




Figure 3: The two inputs for one instance: agent-to-goal distances D and goal-to-goal distances G , both normalized BFS distances with walls respected.

Tokenization. Each scalar D_{ij} is projected independently to a d -dimensional vector by a shared $\text{Linear}(1 \rightarrow d)$, producing an $N \times M$ grid of tokens. Because the projection is shared, the parameter count does not depend on N or M .

Goal context. Each scalar G_{jk} is embedded and the results are summed over k , giving one context vector per goal j that is added to that goal’s column of tokens. The sum-pooling step is what matters here: it maps a variable-length set of goal-to-goal distances onto a fixed-width vector, so M can be anything.

Attention. Each block applies row attention, in which an agent attends over its own goals and learns how they relate to each other, then column attention, in which a goal attends over the candidate agents and learns which is best placed to take it, then a feed-forward layer. Blocks are stacked L deep. Attention operates on variable-length sequences by construction.

No positional embeddings. Agents and goals are unordered sets in an mTSP, so there is no meaningful “agent 1”. Leaving out positional embeddings makes the network permutation-equivariant in both axes and removes the last size-dependent component, which is why the same weights can process a 2×2 instance and a 50×100 one.

Output. A shared $\text{Linear}(d \rightarrow 1)$ reduces each token to a single logit, giving an $N \times M$ logit matrix. A softmax down each column yields P , where P_{ij} is the probability that agent i visits goal j and $\sum_i P_{ij} = 1$. This is exactly the structure Section 2.2 requires. Taking the arg-max of each column decodes the binary allocation.

Mixed-size training. Instances of different shapes cannot be stacked into one tensor, so a `MixedSizeBatchSampler` builds shape-homogeneous batches. Each batch holds a single (N, M) while consecutive batches differ, which trains one model jointly across many configurations without padding or masking.




Figure 4: The size-agnostic path from an instance to an allocation. D and G become an $N \times M$ token grid, row/column attention blocks process it, a shared head produces $N \times M$ logits, a column softmax gives a per-goal distribution over agents, and arg-max decodes the allocation. No stage depends on the values of N or M .

3.4 Loss

Training minimizes

$$\mathcal{L} = \mathcal{L}_{\text{CE}} + \lambda \mathcal{L}_{\text{MinSum}}, \quad \mathcal{L}_{\text{CE}} = -\frac{1}{M} \sum_j \sum_i Y_{ij} \log P_{ij}, \quad \mathcal{L}_{\text{MinSum}} = \frac{1}{M} \sum_{i,j} P_{ij} D_{ij},$$

with $\lambda = 0.1$. The cross-entropy term is per-goal imitation of the solver and does most of the work. The min-sum term is a mild geometric regularizer that penalizes assigning distant goals and prevents degenerate near-uniform predictions early in training. It is only a first-hop proxy, with no tour chaining and no collision term, so the true mTSP objective reaches training solely through the labels Y .

3.5 Data generation and deployment

Each sample is produced by placing agents and goals on a map, computing D and G by BFS, and calling the solver for the optimal Y . Instances with unreachable goals are rejected, as are those that exceed a CBS node budget of 50k expansions. That second filter addresses a failure mode found at high wall density, where an instance passes the reachability check but CBS escalates constraints without ever terminating.

At deployment the arg-max allocation is decoded, each agent’s visit order is fixed by brute-force enumeration over its assigned goals, and CBS plans the paths. When an allocation admits no collision-free plan, the next candidate in a probability-ranked chain is tried, taken best-first over the joint probability $\prod_j P_{a_j,j}$ up to $k = 3$. This is the learned counterpart of the solver’s k -best escape.

4 Experimental Setup and Metrics

Experiments run on procedurally generated grids and on the four 32×32 benchmark maps that RobustMCPF itself uses (Figure 5). Only **maze** and **room** have genuine corridor and room structure, while **empty** and **random** are unstructured. That distinction drives several of the results below.




Figure 5: The four 32×32 benchmark maps: **empty**, **random-32-32-20**, **maze-32-32-2** and **room-32-32-4**. Only the last two are structured.

Allocation-quality metrics compare predicted and ground-truth allocations directly, with no planning involved. Per-goal accuracy is the fraction of goals routed to the solver’s chosen agent. It is a lenient measure, since one wrong goal out of M still scores $(M-1)/M$. Full-assignment accuracy is the all-or-nothing indicator that every column matches, and it empirically tracks

(per-goal accuracy) ^{M} . A proxy cost ratio compares only the first-hop BFS cost of the two allocations. Because it ignores tour chaining and collisions, values below 1.0 reflect the metric’s mismatch with the true objective and not the network beating the solver.

Execution-cost metrics run both pipelines end to end, where a plan’s cost is the total collision-free path length in grid steps that CBS returns. Cost ratio is the per-instance c_{nn}/c_{sol} averaged over instances, so 1.0 means the network matched the exact solver. Exact match is the fraction of instances whose two integer costs are identical, and it runs far higher than full-assignment accuracy because many distinct allocations are cost-equivalent ties. Diff is the mean gap in grid steps, speedup is the ratio of mean single-instance wall-times on the same node, and the infeasible and fallback counters report how often the candidate chain was needed. These are the numbers that matter, because they measure the plan an agent would actually execute.

5 Design Study: What Shaped the System

Fifteen experiments preceded the final system. Instead of recounting them chronologically, we report the five findings that determined the design, each of which fixes one property described in Section 3. Table 1 collects the supporting numbers.

D1. Attention wins, and increasingly so with scale. Against an MLP and a DeepSets encoder on identical data, the row/column transformer led on every metric, and its margin over the MLP grew from 1.7 points of full-assignment accuracy at $N=2$ to 6.7 points at $N=5$. Cross-goal interdependence is what widens as problems grow, and attention is what captures it. This fixes the backbone.

D2. Representation dominates everything else. Adding G lifted full-assignment accuracy by 14 to 18 points at every agent count. For comparison, doubling the training set from 10k to 20k samples bought 0.4 points, and every capacity increase we tried bought at most 1.4, with the largest configuration diverging outright. D on its own cannot express the fact that two goals are adjacent and should be bundled. G supplies precisely that, and it outweighs architecture, data and capacity combined by more than tenfold. This fixes the input.



D3. One size-agnostic model beats a family of specialists. A universal model trained jointly across fifteen (N, M) configurations matched or beat the per-size specialists on all but the easiest configuration, gaining 2.4 to 3.5 points. It also transferred upward. Evaluated on $N=5$, which it had never seen in training, it beat the specialist trained on $N=5$ by 3.5 points. Cross-configuration training acts as a regularizer and pushes the model toward the allocation principle instead of size-specific patterns. This justifies the universal architecture.

D4. Allocation accuracy is the wrong yardstick. Measured end to end, the fraction of instances whose plan costs exactly what the solver’s costs (64–97%) far exceeds the fraction on which the same allocation is chosen (44–91%). Many disagreements are ties, with two agents equidistant from a goal and either choice equally good, so a model judged on allocation accuracy alone looks considerably worse than it performs. This is why every headline number in this report is an execution cost.

D5. Geometry transfers, scale does not. A model trained only on small random grids and evaluated zero-shot on the four real maps at small N, M reached a cost ratio of 1.019, meaning no meaningful degradation. Structured maps even proved easier than open ones, because walls remove the ties that an open grid leaves to chance. Pushed to large N, M , the same model collapsed to 1.246 against a map-trained model’s 1.048, with a revealing signature: at $M=30$, adding agents hurt the model trained on $N \leq 5$ while helping the in-range model (Figure 6b). Separately, training on cheap random grids matched training on the real maps everywhere except the maze, where random obstacles never reproduce corridor structure. This is why the final models are retrained at the target scale and on the target maps, and it is the direct motivation for Section 6.

Table 1: Design decisions and the evidence behind them. Effects are on full-assignment accuracy (D1–D3) or execution-cost ratio (D4–D5).

Decision	Effect	Evidence
Transformer backbone	+1.7 → +6.7 pt	vs. MLP, margin growing over $N=2 \dots 5$
Add goal-goal distances G	+14.2 → +18.0 pt	with vs. without G , all agent counts gains already flattening at 20k
More data (10k → 20k)	+0.4 pt	largest configuration diverges at lr 10^{-3}
More capacity (to h256/L6)	$\leq +1.4$ pt	vs. per-size specialists; +3.5 pt zero-shot at $N=5$
Size-agnostic universal model	+2.4 → +3.5 pt	
Judge by CBS execution cost	64–97% vs. 44–91%	exact-cost match vs. exact-allocation match
Train on the target map family	1.048 vs. 1.075	fixed-map vs. random; the gap is entirely the maze (1.029 vs. 1.127)
Retrain at the target scale	1.048 vs. 1.246	map-trained vs. small-scale model at large N, M
Smaller network in-distribution	1.048 vs. 1.052	h128/L6 vs. h256/L8; $2.5\times$ the parameters bought nothing

(a) the G ablation

(b) the scale-transfer scissors

Figure 6: The two findings that set up the final study. (a) Goal-to-goal distances lift full-assignment accuracy by 14 to 18 points at every agent count, the dominant design decision (D2). (b) At $M=30$, adding agents hurts the small-scale-trained model but helps the map-trained one, so generalization fails across problem scale, not across map geometry (D5).

Why this leads to the final study

D5 showed that a model has to be trained on structured geometry in order to handle it, since the entire advantage of map-based training over random-grid training came from the maze. If geometry-specific structure matters that much, the natural next question is whether a dedicated model per map outperforms one generalist serving all four. That is what Section 6 tests.

6 Final Study: Scale Limits and Specialists versus a Generalist

6.1 Research questions

- RQ1.** How far can N and M grow before the exact solver becomes impractical, on the order of a minute per instance, and what does that imply for what can be labeled and measured?
- RQ2.** For each model role, which training distribution works better: many instances at regular size, or fewer instances at large size?
- RQ3.** Do per-map specialists outperform a single generalist, in quality and in wall-time, by enough to justify routing an instance to a map-specific model?

6.2 Design

The scale wall. We sweep N and M on each map until the solver’s mean wall-time approaches 60 seconds. This bounds the study in three ways that should not be conflated. The labeling ceiling is the largest (N, M) at which the solver can still produce training labels in acceptable time, which bounds what the large-instance dataset can contain. The evaluation ceiling is the largest (N, M) at which a solver baseline exists at all; past that point the cost ratio has no denominator and cannot be computed. Beyond both, only solver-independent quantities remain measurable.

The model matrix. Five roles, namely one generalist trained on all four maps and four specialists trained on `empty`, `random`, `maze` and `room` respectively, are each trained on two datasets: (A) many regular-size instances, and (B) fewer large instances reaching toward the scale wall. This gives ten models, all h128/L6.

Why only the smaller architecture. We fix the network to h128/L6 because it beat the larger h256/L8 in-distribution on every aggregate metric, with a 1.048 versus 1.052 execution-cost ratio and 0.433 versus 0.402 exact match, while running roughly 20% faster. The extra $2.5\times$ in parameters bought no accuracy, and the larger model reached its validation optimum early and then overfitted. The ranking does reverse far out of distribution, where the larger model extrapolates better, but dataset B places the large-instance regime inside the training distribution by construction, so the in-distribution verdict is the applicable one.

TODO [TEXT — 1–2 sentences]

Replace the justification above if the actual reason for restricting to the smaller model differed. The text reconstructs the rationale from the earlier h128-vs-h256 comparison; confirm it, or substitute the real reasoning.

Protocol. All models see identical instances for a given (map, N , M) cell, which allows paired comparison. The dataset choice for RQ2 is made on validation data, and the specialist-versus-generalist comparison for RQ3 is reported on held-out test data, so the selection step cannot inflate the final result. The generalist is scored per map and not pooled. Without that, a uniformly mediocre generalist would be indistinguishable from one that is strong on three maps and weak on the fourth.

TODO [TABLE — T2, the model inventory]

Ten rows: model id, role, training map(s), dataset (A/B), N and M ranges, samples per configuration and in total, epochs, learning rate, batch size, gradient clipping, seeds, parameter count, best validation loss and epoch, training wall-time. Note any model that hit its wall-clock limit before finishing its epoch budget, since a wall-limited model must not be presented as a converged one. *Source: `model_inventory.csv` (data request §6.6).*

6.3 Results

TODO [FIGURE + TEXT — R1, the solver scale wall]

Figure: exact-solver wall-time against problem size, one line per map, horizontal marker at 60s, log scale on time. *Text (~100 words):* where the wall falls on each map, whether it is driven by N or M , and the three ceilings stated numerically. *Source: `solver_scale_wall.csv` (data request §3).*

TODO [TABLE + FIGURE + TEXT — R2, dataset selection (RQ2)]

Five head-to-heads, one per role, each between its dataset-A and dataset-B model on a common (N, M) grid spanning both regimes. *Figure*: cost ratio against problem size, one line per dataset, showing where the curves cross. *Table*: winner per role with the margin. *Text* (~ 150 words): whether the crossover is consistent across maps, and whether many-small or few-large is the better use of a fixed labeling budget. *Source*: *stage1_dataset_selection/* (data request §4).

TODO [TABLE + FIGURE + TEXT — R3, specialists versus the generalist (RQ3)]

The five selected models on all four maps. *Table*: per map, specialist and generalist side by side on execution-cost ratio, exact match, mean diff, per-goal and full-assignment accuracy, and mean wall-time. *Figure*: grouped bars per map, specialist against generalist, on cost ratio and on speed. *Text* (~ 200 words): whether the specialist advantage is real, whether it is uniform across maps or concentrated on the structured ones (the D5 prediction), and whether it holds on timing as well as quality. *Source*: *stage2_expert_vs_general/* (data request §5).

TODO [TABLE + TEXT — R4, the cost of misrouting]

The full 5×4 matrix of every model on every map. The diagonal answers whether specialists are better, and the off-diagonal answers what it costs when a router picks wrong. *Text* (~ 120 words): expected gain under a perfect router, the misrouting penalty, and therefore how confident map identification has to be for routing to pay. *Source*: *stage2_expert_vs_general/* off-diagonal cells.

6.4 Hypotheses

These predictions are recorded before the results are available, so that they can be confirmed or refuted rather than rationalized afterward.

On RQ2, we expect dataset A to win at small (N, M) and dataset B at large (N, M) , with a crossover somewhere between, since each dataset places its own regime in-distribution. The interesting quantity is not which one wins overall but where the crossover sits, because that determines how a fixed labeling budget should be spent.

On RQ3, we expect the specialist advantage to be largest on **maze** and smallest on **empty** and **random**, following D5. Training data mattered only for structured geometry, so specialization should pay only where there is structure to specialize on. If instead the specialists win uniformly across all four maps, the explanation would be capacity, with one network split across four distributions, rather than geometry. That would be a more interesting result than the one we predict.

On timing, we expect specialists and generalist to be indistinguishable, since the architecture is identical and a forward pass costs the same either way. Any measured difference should be attributed to the environment rather than the model, and one that survives scrutiny would be worth investigating.

TODO [TEXT — ~ 150 words]

After the results land, add a short verdict subsection: which hypotheses held, which failed, and what the failures imply. Keep the predictions above unchanged, since a refuted prediction that is honestly reported is worth more than one quietly edited to match.

7 Discussion

When the learned allocator is the right choice. It replaces the allocation step only, so it pays off where allocation dominates and where many instances have to be solved. At large N and M , LKH’s TSP grows expensive while the forward pass stays close to flat, which yields 5.9–9.4× end-to-end speedup at roughly 20% above optimal cost. Under batched, high-throughput use, allocation-only inference is orders of magnitude faster than the solver. On unstructured geometry, even cheap random-grid training transfers, so no map-specific data collection is needed. The opposite cases are equally clear. On small instances the exact solver already answers in milliseconds and CBS dominates both pipelines, leaving nothing to save. Where exactness is contractually required, an approximate allocator is inadmissible whatever its speed. And on structured maps the network has to have trained on that structure, so an unseen structured environment is a poor fit.

Limitations. Goal-visit ordering is still brute-force over each agent’s assigned goals, which is $O(k!)$ and becomes the pipeline bottleneck at seven or eight goals per agent. The allocation itself is sub-millisecond, but the classical step around it is not. CBS execution cost is not differentiable, so training optimizes an imitation proxy and not the objective actually being measured. Imitation also caps quality at the teacher’s level, meaning the network can match the exact solver but never beat it. Finally, the evaluation rests on one 32×32 benchmark family under basic, deterministic MAPF. Orientation, action delays, and larger or non-grid environments are untested.

8 Conclusion

A compact transformer reading two distance matrices reproduces an exact solver’s goal allocation closely enough that the resulting collision-free plans cost 1–6% more than optimal in-distribution, with a single set of weights serving any number of agents and goals. The decisive design choice turned out to be representation, not architecture or scale, since goal-to-goal distances were worth more than every other improvement combined. The model generalizes across map geometry but not across problem scale, and its speed advantage appears in exactly the regimes where the exact solver’s does not, at scale and in batch.

TODO [TEXT — 2–3 sentences]

Add the final study’s conclusion: whether map specialists beat the generalist, whether routing is justified, and which training distribution proved the better use of a labeling budget.

Future work. The solver’s true tour cost is already computed during data generation and then discarded. Storing it would enable a tour-aware loss term $\sum_{j,k} (\sum_i P_{ij} P_{ik}) G_{jk}$ that directly penalizes splitting adjacent goals, targeting the failure mode that dominates at high M . A learned goal-ordering head would remove the last classical bottleneck. Approximate conflict penalties or policy-gradient signals could push training toward the true MAPF objective instead of the imitation proxy.

Acknowledgements

The RobustMCPF solver, the exact LKH and CBS pipeline used here for ground truth and as the comparison baseline, is the work of Yehonatan Kidushim, Avraham Natan, Roni Stern and Meir Kalech, and is not a contribution of this project. We thank the authors for permission to use it.

TODO [TEXT — verify before submission]

The references below are standard works, but their venues, years and page numbers have not been checked against the originals. Verify each one. Related work (Section 2.4) is deliberately one paragraph; expand it into a standalone section if the grader expects a fuller treatment.

References

- [1] Y. Kidushim, A. Natan, R. Stern, and M. Kalech. Robust Multiagent Combinatorial Path Finding. *AAAI*, 2026.
- [2] G. Sharon, R. Stern, A. Felner, and N. Sturtevant. Conflict-based search for optimal multi-agent pathfinding. *Artificial Intelligence*, 219:40–66, 2015.
- [3] K. Helsgaun. An effective implementation of the Lin–Kernighan traveling salesman heuristic. *European Journal of Operational Research*, 126(1):106–130, 2000.
- [4] R. Stern et al. Multi-agent pathfinding: Definitions, variants, and benchmarks. *Symposium on Combinatorial Search (SoCS)*, 2019.
- [5] N. Sturtevant. Benchmarks for grid-based pathfinding. *IEEE Transactions on Computational Intelligence and AI in Games*, 4(2):144–148, 2012.
- [6] H. Ma and S. Koenig. Optimal target assignment and path finding for teams of agents. *AAMAS*, 2016.
- [7] O. Kaduri, E. Boyarski, and R. Stern. Algorithm selection for optimal multi-agent pathfinding. *ICAPS*, 2020.
- [8] G. Sartoretti et al. PRIMAL: Pathfinding via reinforcement and imitation multi-agent learning. *IEEE Robotics and Automation Letters*, 4(3):2378–2385, 2019.
- [9] O. Vinyals, M. Fortunato, and N. Jaitly. Pointer networks. *NeurIPS*, 2015.
- [10] I. Bello, H. Pham, Q. V. Le, M. Norouzi, and S. Bengio. Neural combinatorial optimization with reinforcement learning. *ICLR Workshop*, 2017.
- [11] W. Kool, H. van Hoof, and M. Welling. Attention, learn to solve routing problems! *ICLR*, 2019.
- [12] Y. Bengio, A. Lodi, and A. Prouvost. Machine learning for combinatorial optimization: a methodological tour d’horizon. *European Journal of Operational Research*, 290(2):405–421, 2021.

A Appendix

TODO [TABLE — full per-configuration results]

Complete per-configuration tables for the final study: every (model, map, N , M) cell with cost ratio, exact match, mean/max/std diff, infeasible and fallback rates, timings, and n . The body carries only aggregates and figures. *Source: aggregate files accompanying stage2.expert-vs-general/.*

TODO [TABLE — reproducibility]

Hyperparameters and seeds for all ten models, cluster node types used for timing, torch version, and whether inference ran on CPU or GPU. Timing numbers are uninterpretable without this.

TODO [FIGURE — optional, training convergence]

Train and validation loss and per-goal accuracy against epoch for the ten models, with the best-validation checkpoint marked. Worth including if the models show the same overfitting signature as earlier runs, where validation loss turned upward well before the epoch budget ended and best-validation checkpointing acted as implicit early stopping. *Source: train_logs/ (data request §6.5).*

TODO [TEXT — raw data pointers]

Record where the per-instance CSVs live (report/data/exp16/), the column schema, and the cluster job identifiers, matching the convention used for earlier experiments.